

UG HW6: Semaphores for xv6

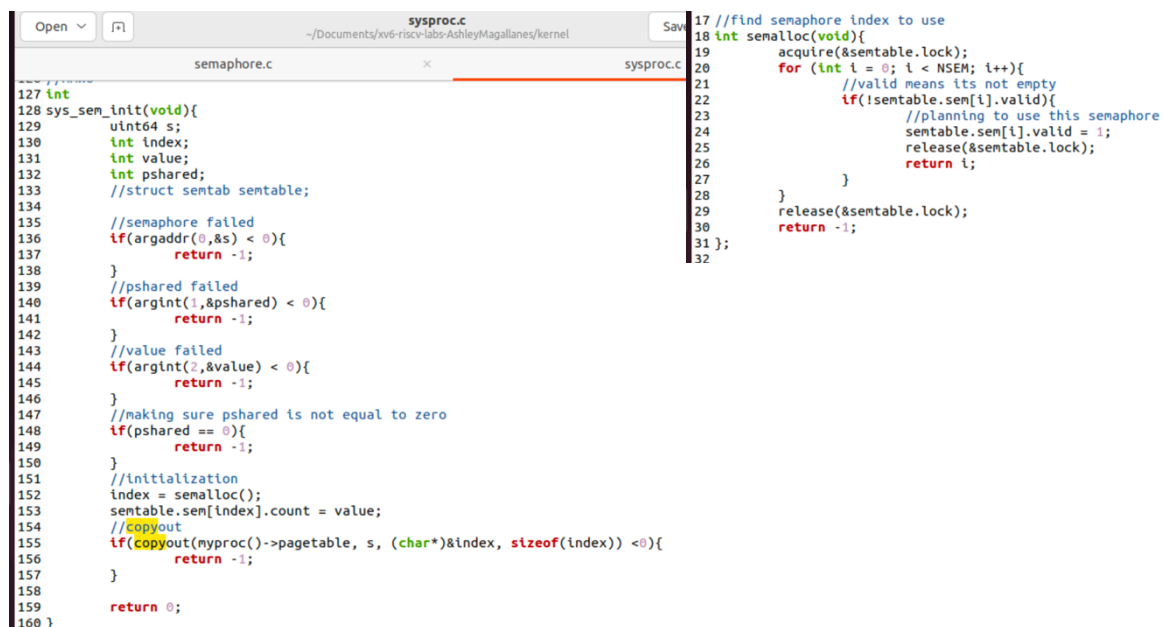
Task 3. Implementation of sem_init(), sem_wait(), sem_post(), and sem_destroy().

The implementation of the semaphore revolves around four main system calls:

1. System call for sem_init()

This method initializes the semaphore given the address passed as an argument. To do this I need to find an unused location. Therefore, here the helper methods that I implemented in the semaphore.c file for task 2 come in handy. One of the helper methods is semalloc(), which as needed it finds a semaphore index to use. This method goes through NSEM or the size defined value and checks whether each index is valid, and once it finds one that is not valid it means that is empty and it can be used, so it returns that index along with removing its valid characteristic since it will now be used. This helper method made it easier to implement the sys_sem_init().

I started by creating the data variables for each of the arguments and also one for the index being found. I then proceeded to extract the argument values for the semaphore address, the value, and the pshared from the registers. I used araddr() for the semaphore address and argint() for the value and pshared. After making sure all values extracted from registers were done properly then I started the initialization of the semaphore. I set the index to the return value of semalloc(). Therefore, having the variable index hold the next available semaphore index, which then I used to initialize the value. I directly accessed the count field of the semaphore by using the extern structure semtable and passed the index found above, then I set this to the value given as argument. Lastly, in order to complete the method I used copyout() to copy the semaphore from the kernel to the user space memory.



```
127 int
128 sys_sem_init(void){
129     uint64 s;
130     int index;
131     int value;
132     int pshared;
133     //struct semtab semtable;
134
135     //semaphore failed
136     if(argaddr(0,&s) < 0){
137         return -1;
138     }
139     //pshared failed
140     if(argint(1,&pshared) < 0){
141         return -1;
142     }
143     //value failed
144     if(argint(2,&value) < 0){
145         return -1;
146     }
147     //making sure pshared is not equal to zero
148     if(pshared == 0){
149         return -1;
150     }
151     //initialization
152     index = semalloc();
153     semtable.sem[index].count = value;
154     //copyout
155     if(copyout(myproc()->pagetable, s, (char*)&index, sizeof(index)) < 0){
156         return -1;
157     }
158
159     return 0;
160 }
```

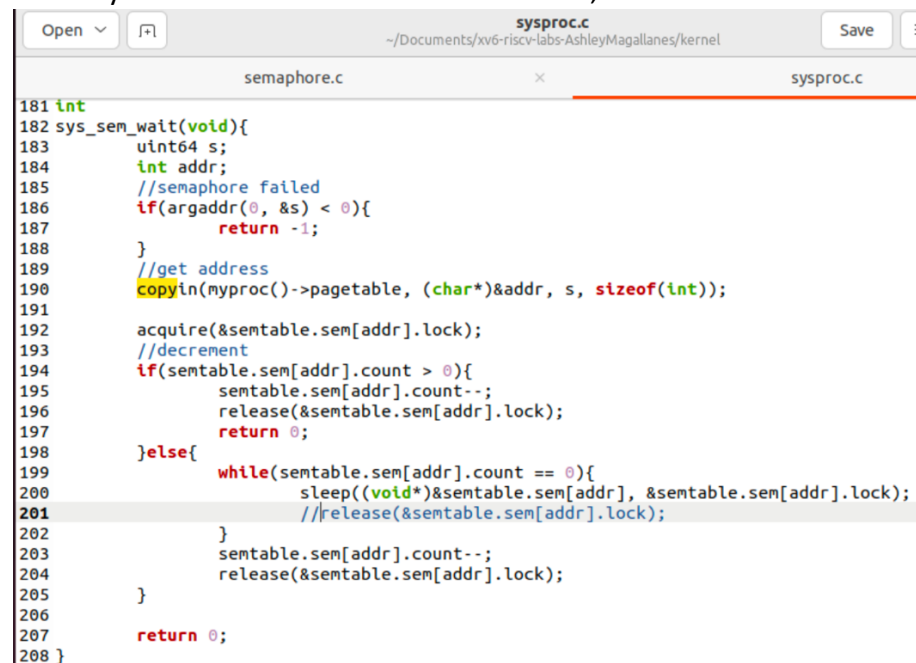
```
17 //find semaphore index to use
18 int semalloc(void){
19     acquire(&semtable.lock);
20     for (int i = 0; i < NSEM; i++){
21         //valid means its not empty
22         if(!semtable.sem[i].valid){
23             //planning to use this semaphore
24             semtable.sem[i].valid = 1;
25             release(&semtable.lock);
26             return i;
27         }
28     }
29     release(&semtable.lock);
30     return -1;
31 }
32 }
```

2. System call for sem_wait()

This method is utilized to decrement or lock the semaphore. The already built function sleep() is useful here since the count value of the semaphore needs to be greater than zero for it to decrement. If the count value of the semaphore is equal to zero then the semaphore needs to be blocked until it is possible to decrement and that is where sleep() plays a role.

I began the implementation of sys_sem_wait() by creating the variables needed, one to hold the argument passed and another for the specific index being used of the semaphore. Then I proceeded to retrieve this argument from the zero register and made sure it didn't fail. Then I used copyin() to get the address or index that would be used to access the count field of the semaphore. In addition, before any decrementing procedure I made sure to have concurrency control by adding a lock for the semaphore at the specific index being accessed.

Then as mentioned I need to decrement; therefore I checked if the count value of the semaphore at the index was greater than zero, and if so then the count was decremented and the lock was released. However, if the value of count is equal to zero then the semaphore was blocked using the sleep function. Then after the count was not zero anymore then the decrement was done, and the lock was released.



```
181 int
182 sys_sem_wait(void){
183     uint64 s;
184     int addr;
185     //semaphore failed
186     if(argaddr(0, &s) < 0){
187         return -1;
188     }
189     //get address
190     copyin(myproc()->pagetable, (char*)&addr, s, sizeof(int));
191
192     acquire(&semtable.sem[addr].lock);
193     //decrement
194     if(semtable.sem[addr].count > 0){
195         semtable.sem[addr].count--;
196         release(&semtable.sem[addr].lock);
197         return 0;
198     }else{
199         while(semtable.sem[addr].count == 0){
200             sleep((void*)&semtable.sem[addr], &semtable.sem[addr].lock);
201             //release(&semtable.sem[addr].lock);
202         }
203         semtable.sem[addr].count--;
204         release(&semtable.sem[addr].lock);
205     }
206
207     return 0;
208 }
```

3. System call for sem_post()

This method is used to increment or unlock the semaphore. It is almost like the opposite of sem_wait(). However, for this method the count value is not required to be greater than zero in order to increment. I started the implementation just like the other two. I retrieved the semaphore address from the arguments passed by accessing the value in register zero. Then again I used copyin() to get the address or index that would be used to access the count field of the semaphore. I also made sure to have concurrency control by adding a lock for the semaphore at the specific index being

accessed. After that I incremented the semaphore count by accessing the external semtable structure and passing in the address to sem. Then I used the wakeup() function passing through the argument as the semaphore index. Lastly, I made sure to just close the lock to avoid any panic arguments.

```

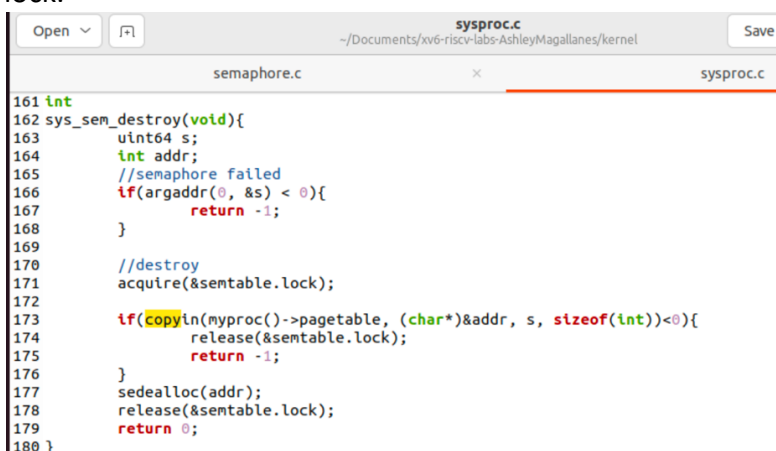
209 int
210 sys_sem_post(void){
211     uint64 s;
212     int addr;
213     //semaphore failed
214     if(argaddr(0, &s) < 0){
215         return -1;
216     }
217     //get address
218     copyin(myproc()->pagetable, (char*)&addr, s, sizeof(int));
219
220     acquire(&semtable.sem[addr].lock);
221     //increment
222     semtable.sem[addr].count++;
223     wakeup((void*)&semtable.sem[addr]);
224
225     release(&semtable.sem[addr].lock);
226
227     return 0;
228 }

```

4. System call for sem_destroy()

This method destroys the semaphore at the address passed by the argument. To do this I need to deallocate the semaphore as well. Therefore, here the helper methods that I implemented in the semaphore.c file for task 2 come in handy. One of the helper methods is sedealloc(), this function invalidates the entry of the given index from the semaphore. The method makes sure that the index passed through the argument is a valid index making sure it is in the range of size of the semtable. Then set the valid field of the semaphore index to zero, meaning that it is not used anymore. This “invalidates” the semaphore.

I implemented the sys_sem_destroy() function by getting the semaphore address of the argument from the register just like in the previous methods. Then I once again made sure to have concurrency control by adding a lock for the semaphore semtable. I used copyin() to get the index that needs to be invalidated. After getting the index I passed it through the argument of the helper method sedealloc(), which already handles the invalidation of this semaphore entry. Lastly, I just released the semtable lock.



```

161 int
162 sys_sem_destroy(void){
163     uint64 s;
164     int addr;
165     //semaphore failed
166     if(argaddr(0, &s) < 0){
167         return -1;
168     }
169
170     //destroy
171     acquire(&semtable.lock);
172
173     if(copyin(myproc()->pagetable, (char*)&addr, s, sizeof(int)) < 0){
174         release(&semtable.lock);
175         return -1;
176     }
177     sedealloc(addr);
178     release(&semtable.lock);
179     return 0;
180 }

```

```

33 //invalidate entry of given index
34 void sdealloc(int index){
35     acquire(&semtable.sem[index].lock);
36     //make sure index arg is valid
37     if(index > -1 && index < NSEM){
38         //invalidate the entry
39         semtable.sem[index].valid = 0;
40     }
41     release(&semtable.sem[index].lock);
42 };

```

The hardest issue that I ran into when completing this task was panic errors in lock releases. I had to move around the locks in order to see where I needed them and where I didn't. One main issue with this was in `sys_sem_wait()` I was trying to release the lock both inside the while and after the whole decrement procedure, therefore I was running into race conditions. However, after some thinking I realized that only leaving the semaphore lock release after the whole procedure would work. Other than that, maybe something else was just understanding how the address accessing worked because at first I did the functions with the whole semaphore, without passing the address and of course this was giving me incorrect and inconsistent output to the testing cases of the `prodcons-sem` command.

Task 4.

Test cases.

| ashley@ashley-virt | ashley@ashley-virt | ashley@ashley-vi | ashley@ashley- |
|--|--|---|---|
| hart 1 starting hart 2 starting init: starting sh \$ prodcons-sem 1 1 producer 5 producing 1 consumer 4 consuming 1 producer 5 producing 2 consumer 4 consuming 2 producer 5 producing 3 consumer 4 consuming 3 producer 5 producing 4 consumer 4 consuming 4 producer 5 producing 5 consumer 4 consuming 5 producer 5 producing 6 consumer 4 consuming 6 producer 5 producing 7 consumer 4 consuming 7 producer 5 producing 8 consumer 4 consuming 8 producer 5 producing 9 consumer 4 consuming 9 producer 5 producing 10 consumer 4 consuming 10 producer 5 producing 11 consumer 4 consuming 11 producer 5 producing 12 consumer 4 consuming 12 producer 5 producing 13 consumer 4 consuming 13 producer 5 producing 14 consumer 4 consuming 14 producer 5 producing 15 consumer 4 consuming 15 producer 5 producing 16 consumer 4 consuming 16 producer 5 producing 17 consumer 4 consuming 17 | \$ prodcons-sem 2 3 producer 10 producing 1 producer 10 producing 2 consumer 8 consuming 1 consumer 7 consuming 2 producer 10 producing 3 consumer 7 consuming 3 producer 10 producing 4 consumer 7 consuming 4 producer 10 producing 5 consumer 7 consuming 5 producer 10 producing 6 consumer 7 consuming 6 producer 10 producing 7 consumer 7 consuming 7 producer 10 producing 8 consumer 7 consuming 8 producer 10 producing 9 consumer 7 consuming 9 producer 10 producing 10 consumer 7 consuming 10 producer 10 producing 11 consumer 7 consuming 11 producer 10 producing 12 consumer 7 consuming 12 producer 10 producing 13 consumer 7 consuming 13 producer 10 producing 14 consumer 7 consuming 14 producer 10 producing 15 consumer 7 consuming 15 producer 10 producing 16 consumer 7 consuming 16 producer 10 producing 17 consumer 7 consuming 17 producer 10 producing 18 consumer 7 consuming 18 producer 10 producing 19 consumer 7 consuming 19 | \$ prodcons-sem 5 2 producer 15 producing 1 producer 15 producing 2 consumer 14 consuming 1 producer 15 producing 3 producer 15 producing 4 producer 15 producing 5 producer 15 producing 6 consumer 13 consuming 2 consumer 13 consuming 3 consumer 14 consuming 4 consumer 14 consuming 5 producer 15 producing 7 producer 15 producing 8 producer 15 producing 9 consumer 14 consuming 6 consumer 13 consuming 7 consumer 14 consuming 8 consumer 13 consuming 9 producer 15 producing 10 consumer 13 consuming 10 producer 15 producing 11 consumer 13 consuming 11 producer 15 producing 12 consumer 13 consuming 12 producer 15 producing 13 producer 15 producing 14 consumer 13 consuming 13 consumer 14 consuming 14 producer 15 producing 15 consumer 13 consuming 15 producer 15 producing 16 consumer 13 consuming 16 producer 15 producing 17 consumer 13 consuming 17 producer 15 producing 18 consumer 13 consuming 18 producer 15 producing 19 consumer 14 consuming 19 | \$ prodcons-sem 3 4 producer 26 producing 1 producer 26 producing 2 consumer 22 consuming 1 consumer 21 consuming 2 producer 25 producing 3 producer 26 producing 4 consumer 22 consuming 3 consumer 21 consuming 4 producer 25 producing 5 consumer 21 consuming 5 producer 25 producing 6 consumer 21 consuming 6 producer 25 producing 7 producer 25 producing 8 producer 25 producing 9 consumer 23 consuming 7 consumer 21 consuming 8 consumer 22 consuming 9 producer 25 producing 10 consumer 21 consuming 10 producer 25 producing 11 consumer 21 consuming 11 producer 25 producing 12 consumer 21 consuming 12 producer 25 producing 13 consumer 21 consuming 13 producer 25 producing 14 consumer 21 consuming 14 producer 25 producing 15 producer 25 producing 16 consumer 22 consuming 15 consumer 21 consuming 16 producer 25 producing 17 consumer 21 consuming 17 producer 25 producing 18 consumer 21 consuming 18 producer 25 producing 19 consumer 21 consuming 19 |
| producer 5 producing 18 consumer 4 consuming 18 producer 5 producing 19 consumer 4 consuming 19 producer 5 producing 20 consumer 4 consuming 20 total = 210 \$ | consumer 7 consuming 19 producer 10 producing 20 consumer 7 consuming 20 total = 210 \$ | producer 15 producing 20 consumer 14 consuming 19 consumer 14 consuming 20 total = 210 \$ | consumer 21 consuming 19 producer 25 producing 20 consumer 21 consuming 20 total = 210 \$ |

The test cases above were done to demonstrate the correct implementation of the `sem_init()`, `sem_wait()`, `sem_post()`, and `sem_destroy()` methods and system calls. I decided to follow three of the output examples provided by Dr. Moore in order to make sure the output was produced properly then I also tried one more of my choice. All tests were correct and concurrent with each other with an output total of 210. Therefore, my implementation of the system calls along with the `semaphore.c` file was properly done.

Kernel bug with our implementation.

I followed the instructions described in the lab therefore, there is a bug in my semaphore implementation. If a user program doesn't deallocate semaphores that were allocated in the kernel semaphore table then a problem would be presented. If the operating system does not clean up after the program, meaning that it doesn't handle these deallocations of semaphores after procedures then the system can fail. If there is too many entries that are not being deallocated then this can become an issue if more processes need availability and there is not enough. This issue of the semaphore table being full of entries can also lead to failing in allocating new entries as well. All of this then can produce unexpected outputs and incorrect cases. However, this issue can be resolved by adding an implementation process to do this deallocation of unused semaphore entries.

Summary:

This homework assignment allowed me to experiment a little more with system calls but more than that I was capable of understanding threads and processes even more. I wasn't sure of why locks were needed but it makes sense that when leading with multiple threads then race conditions can happen and this leads to a shutdown of the system along with panic outputs. I believe that semaphores are important factors in the operating system, and being able to build it from scratch was really interesting!